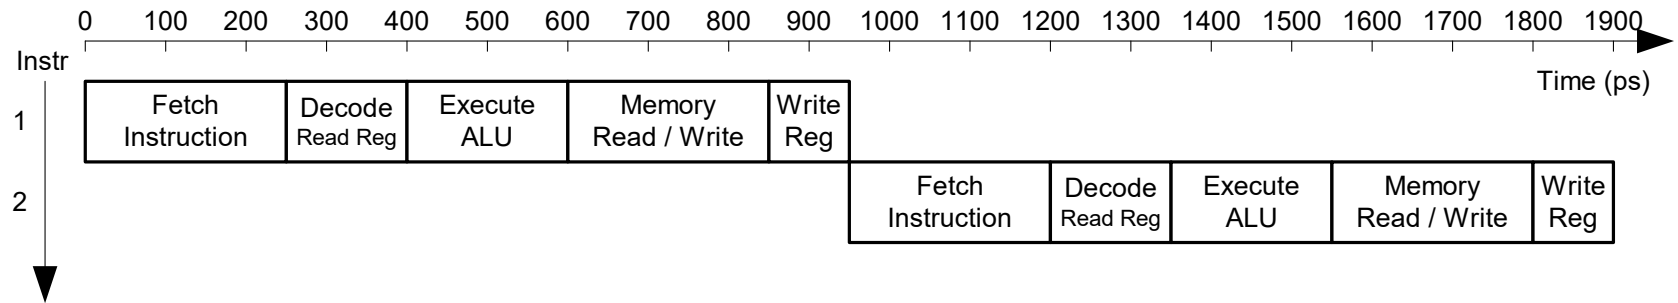


# Pipelined MIPS Processor

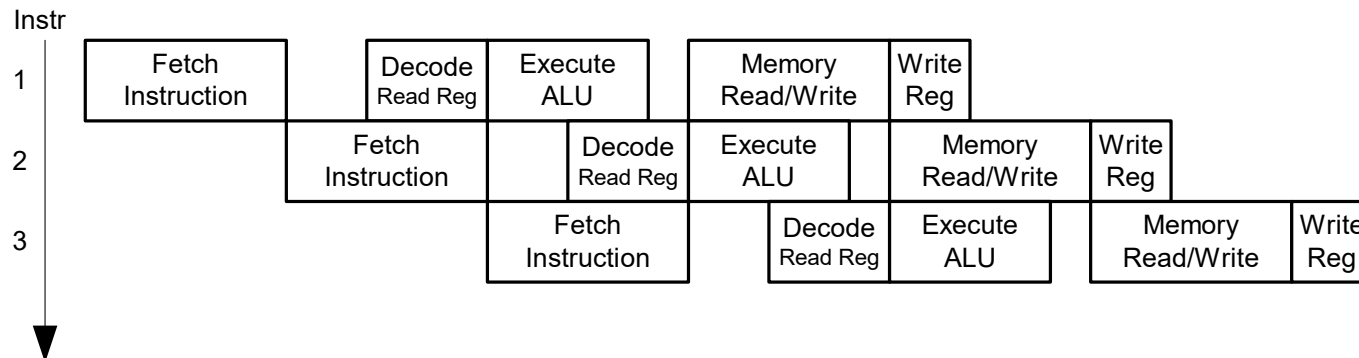
- Temporal parallelism
- Divide single-cycle processor into 5 stages:
  - Fetch
  - Decode
  - Execute
  - Memory
  - Writeback
- Add pipeline registers between stages

# Single-Cycle vs. Pipelined

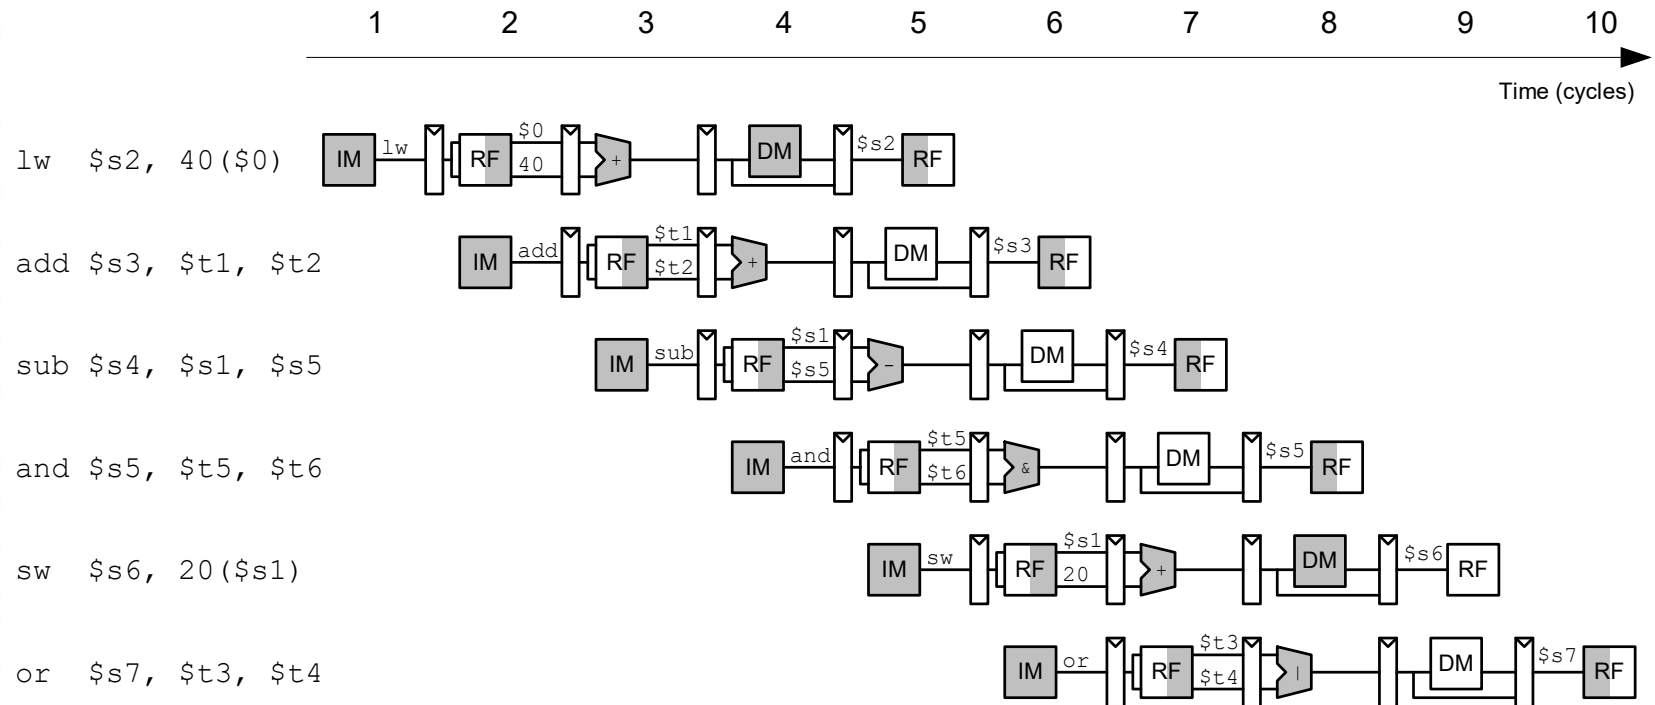
## Single-Cycle



## Pipelined



# Pipelined Processor Abstraction



# Pipelined & RTL



图 4-2 状态机的简单模型

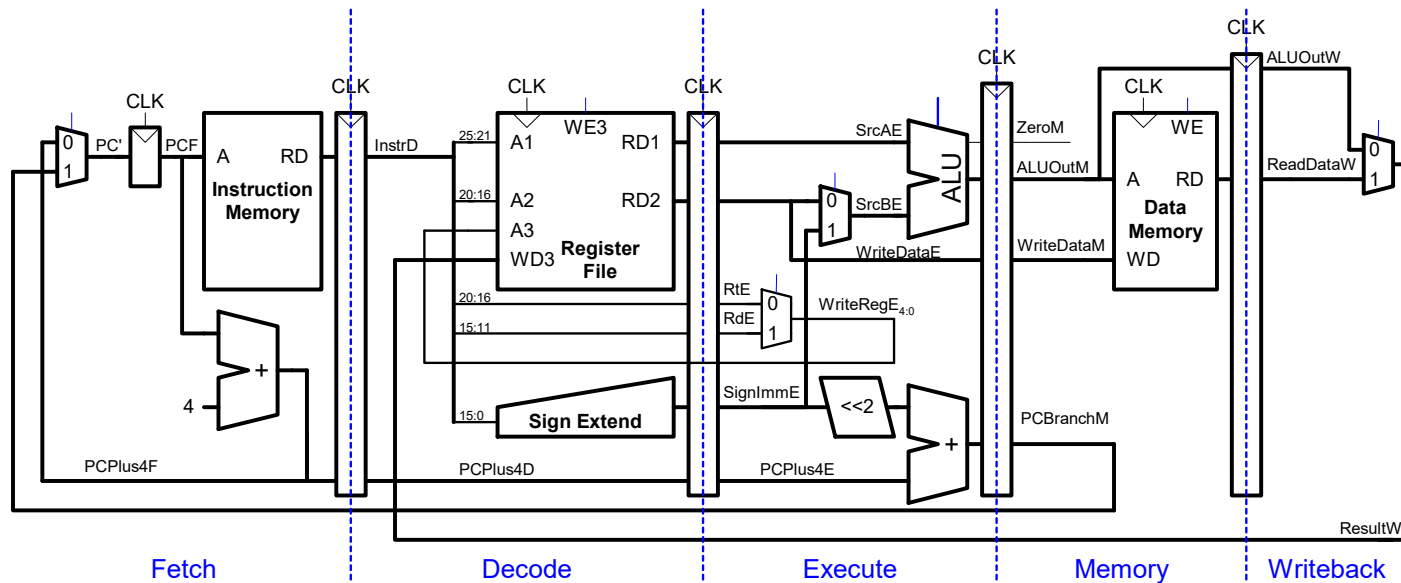
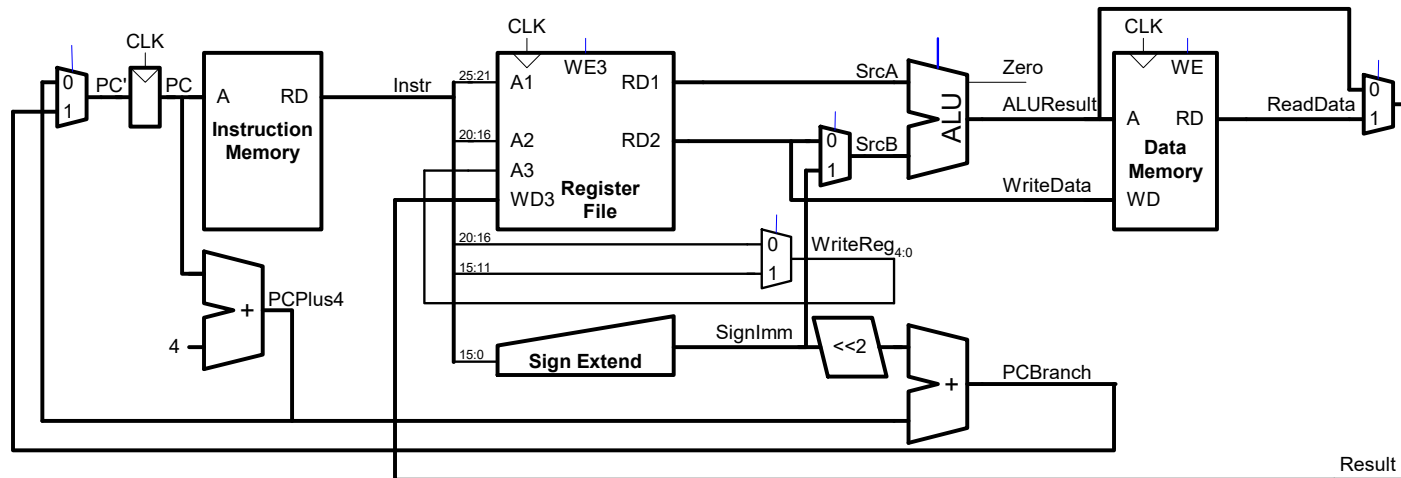


图 4-3 流水线的简单模型 [blog.csdn.net/leishangwen](http://blog.csdn.net/leishangwen)

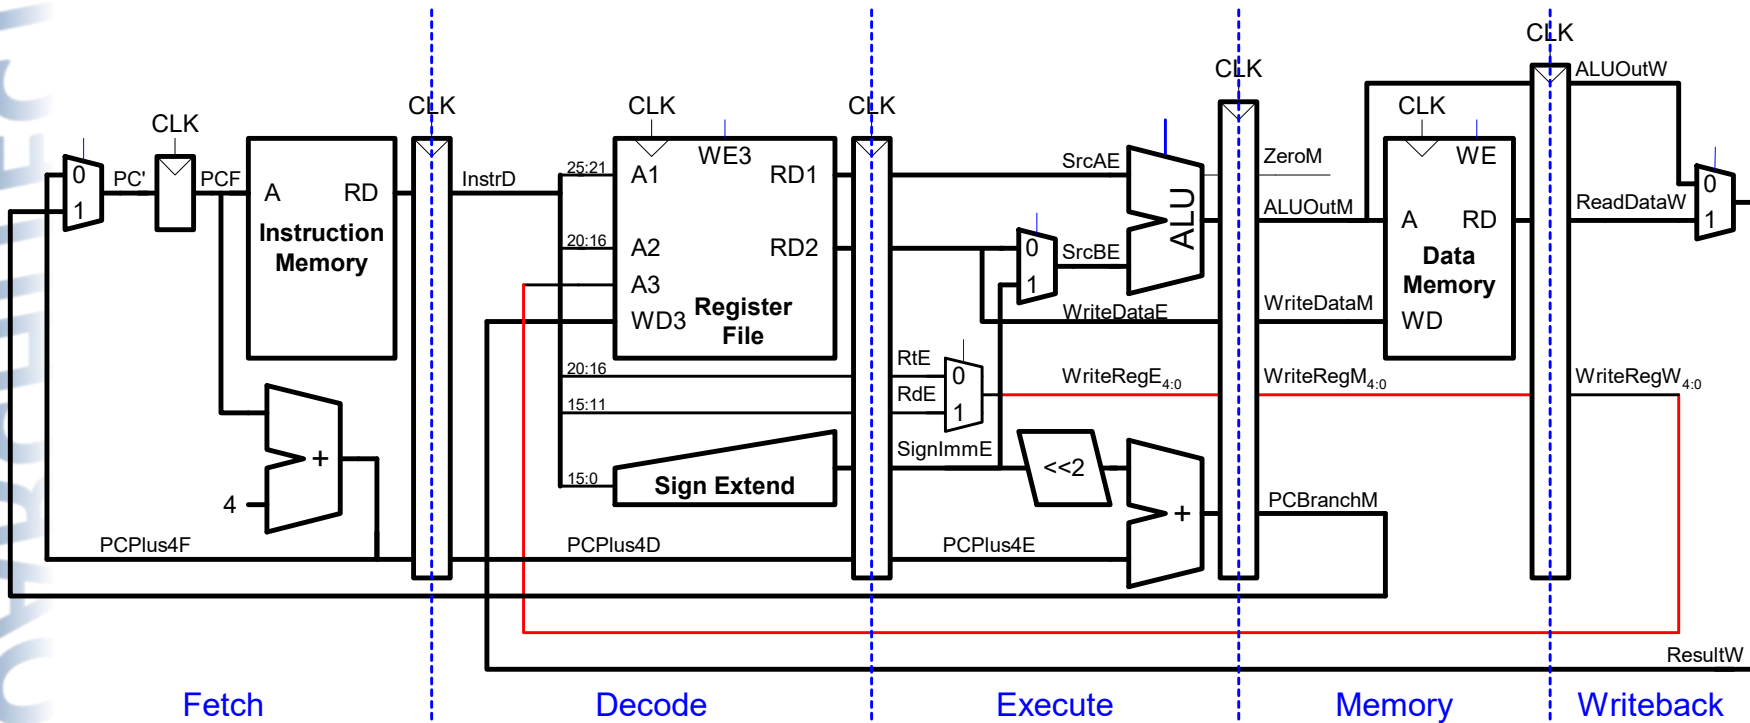
## RTL: Register Transfer Level

The signal translate between registers, and every stage occurs their change for combined logic

# Single-Cycle & Pipelined Datapath

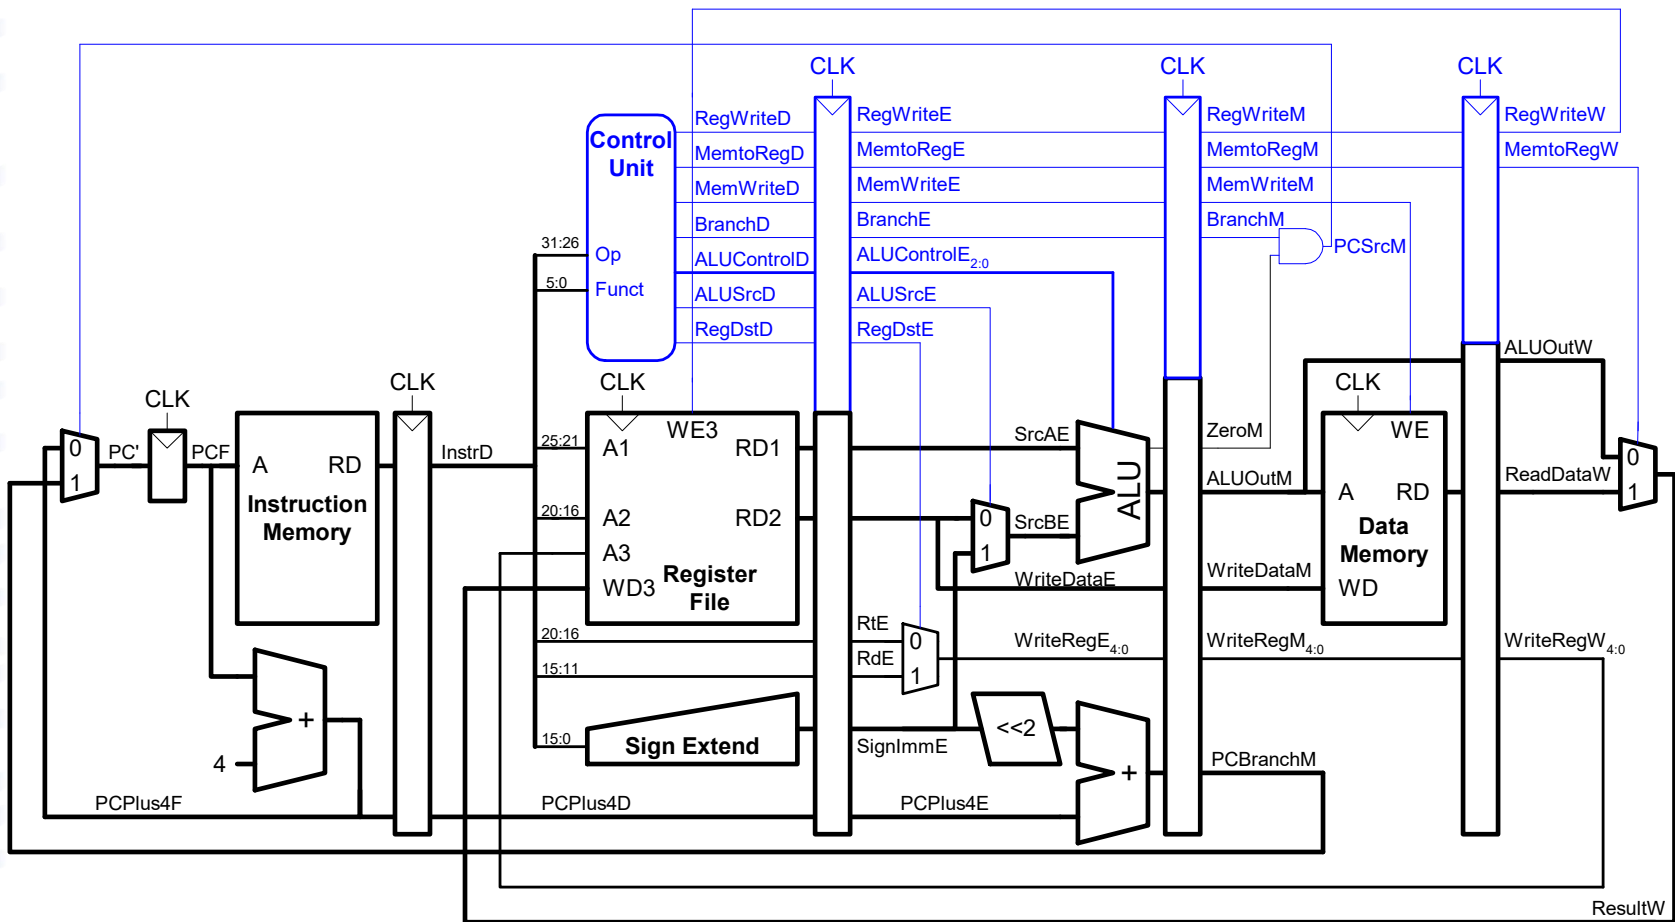


# Corrected Pipelined Datapath



*WriteReg must arrive at same time as Result*

# Pipelined Processor Control



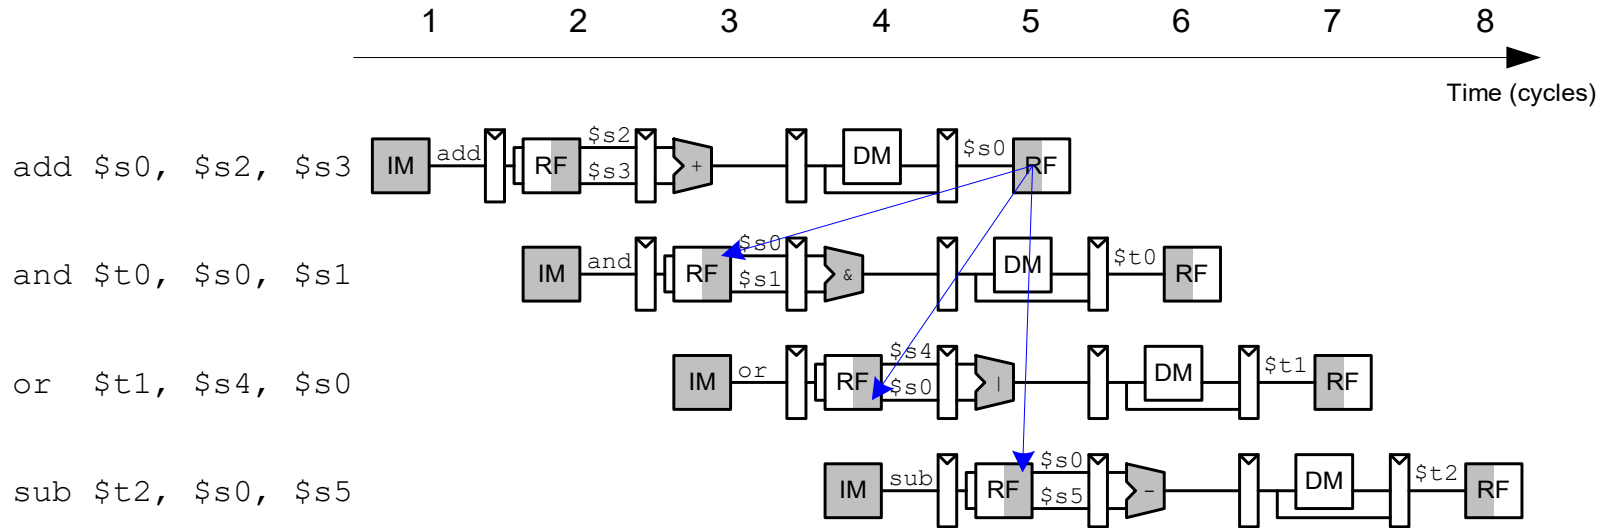
- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage

# Pipeline Hazards

- When an instruction depends on result from instruction that hasn't completed
- Types:
  - **Structural hazard:** We have already solve it by using Harvard architecture. (depart the instruction memory and data memory into two pieces)
  - **Data hazard:** register value not yet written back to register file
  - **Control hazard:** next instruction not decided yet (caused by branches)



# Data Hazard

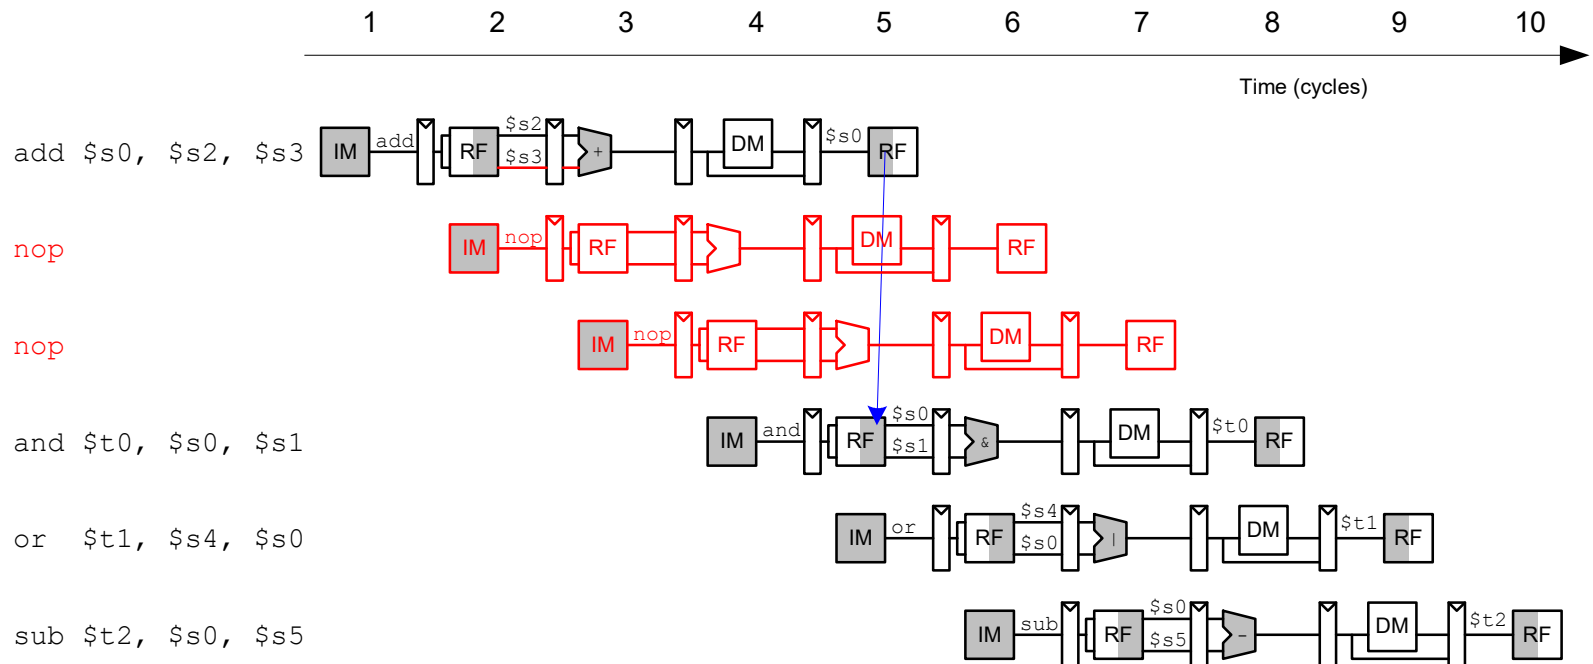


# Handling Data Hazards

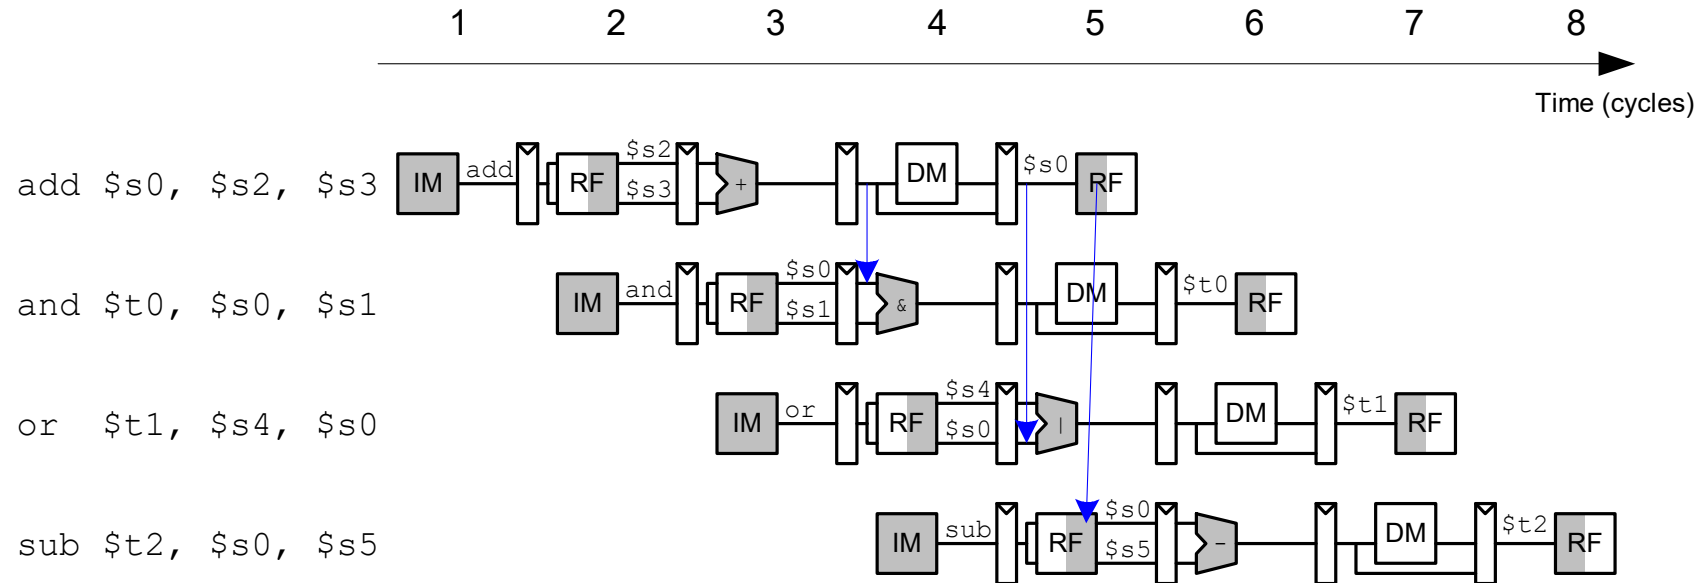
- Insert `nops` in code at compile time
- Rearrange code at compile time
- Forward data at run time
- Stall the processor at run time

# Compile-Time Hazard Elimination


- Insert enough nops for result to be ready
- Or move independent useful instructions forward



# Data Forwarding



# Data Forwarding



Hazard Unit

# Data Forwarding

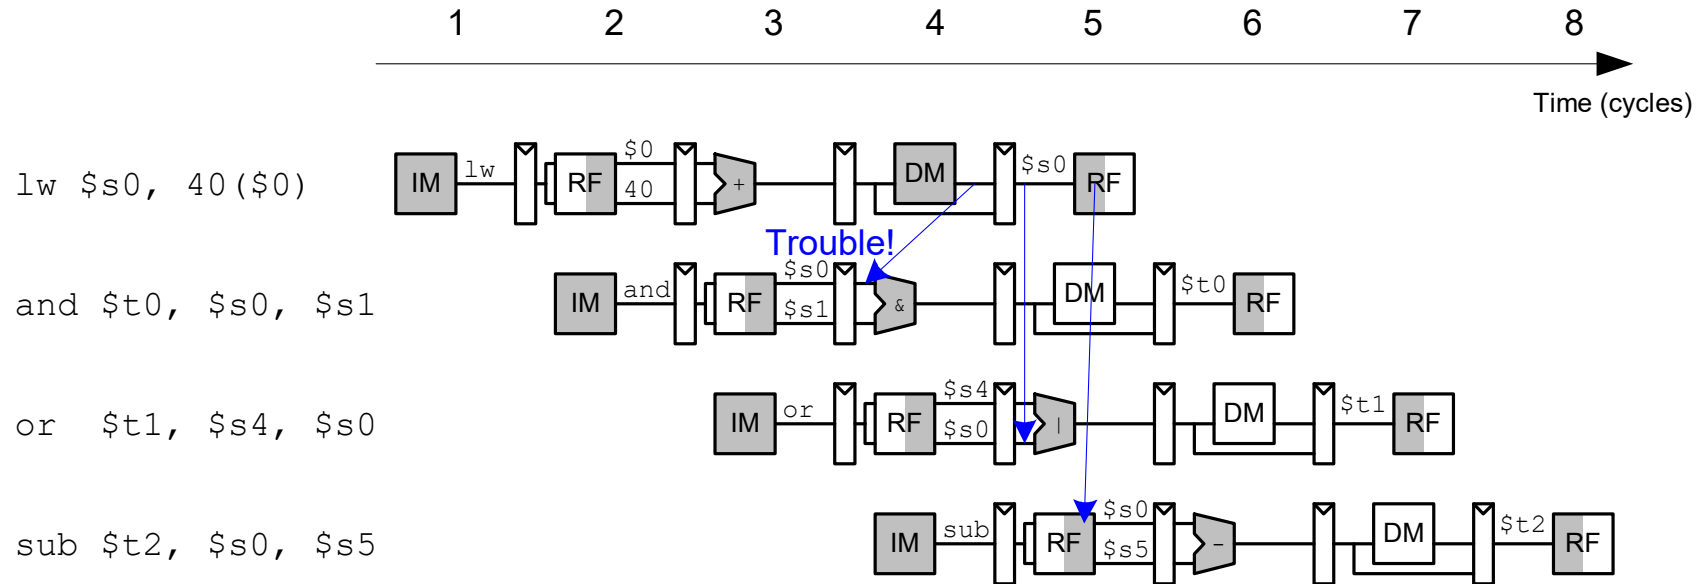
- Forward to Execute stage from either:
  - Memory stage or
  - Writeback stage
- Forwarding logic for *ForwardAE*:

```
if      ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM)
    then ForwardAE = 10
else if ((rsE != 0) AND (rsE == WriteRegW) AND RegWriteW)
    then ForwardAE = 01
else    ForwardAE = 00
```

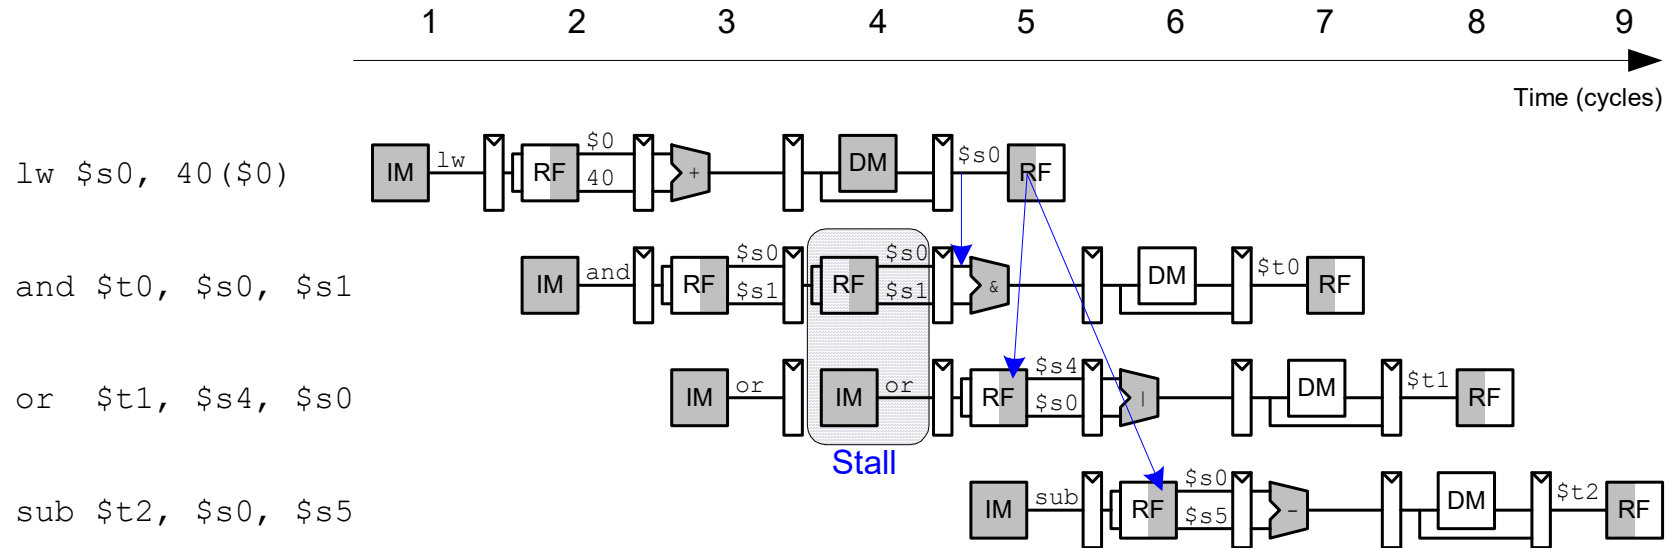
Forwarding logic for *ForwardBE* same, but replace *rsE* with *rtE*



# Stalling




# Stalling





# Stalling Hardware



# Stalling Logic

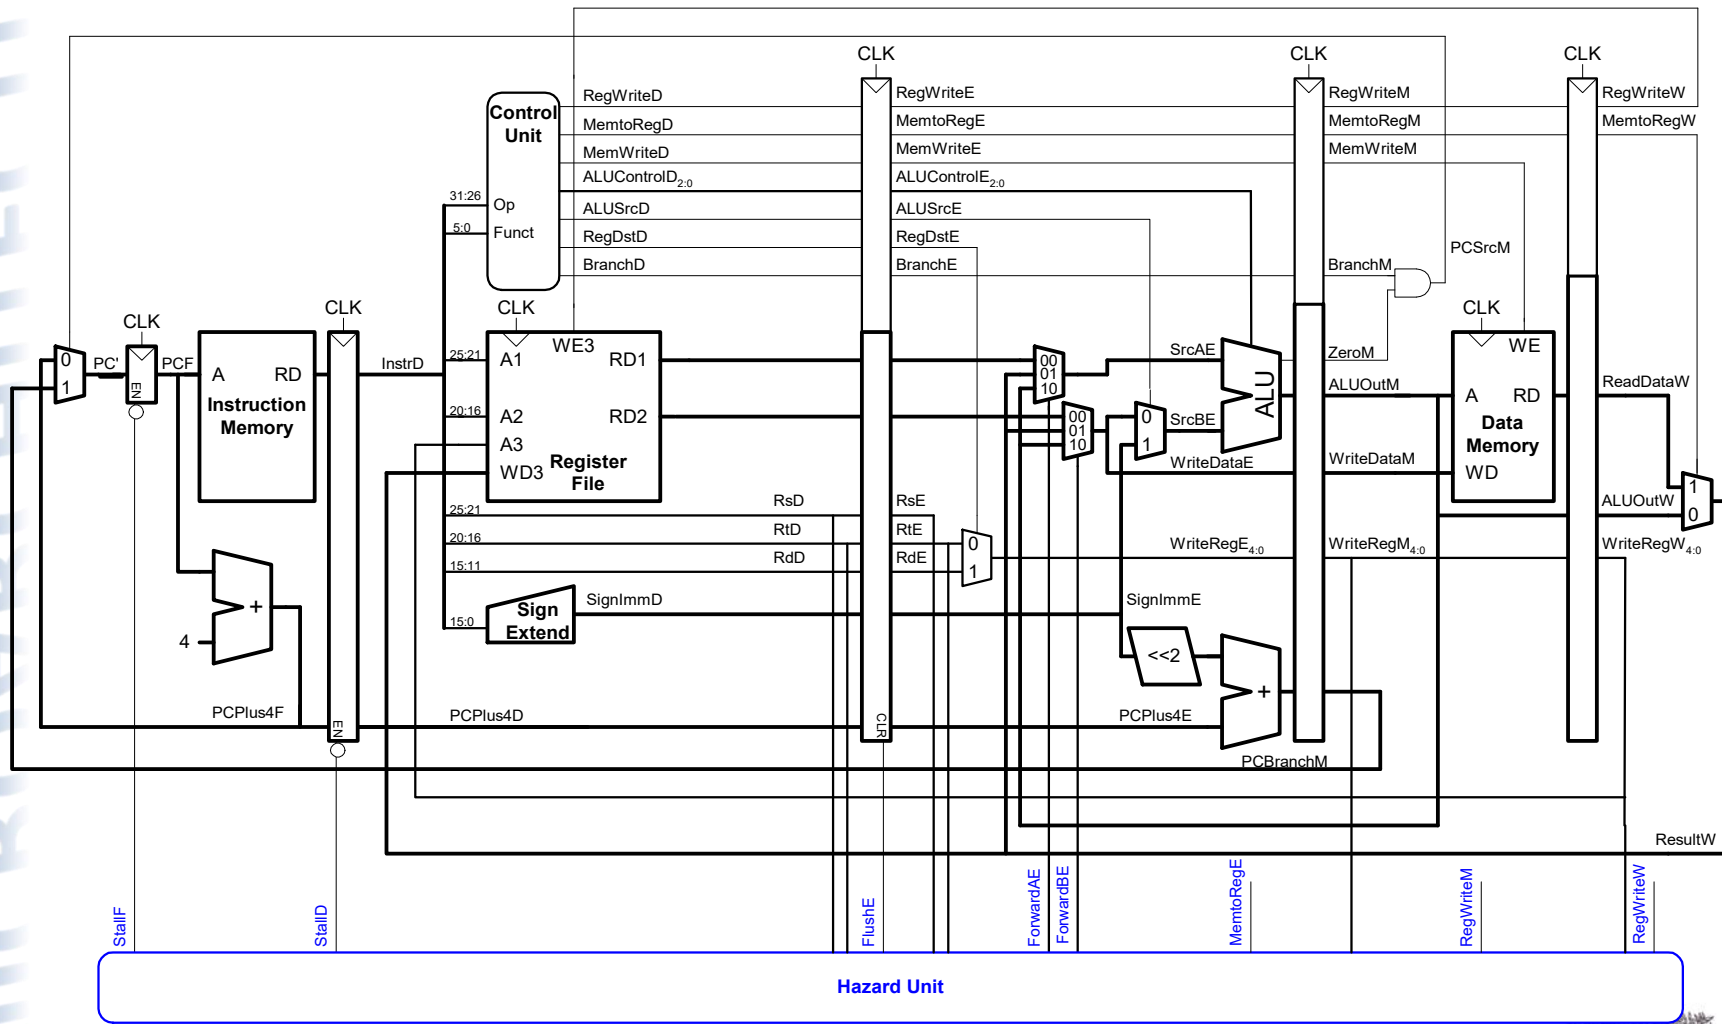
$lwstall =$   
 $((rsD == rtE) \text{ OR } (rtD == rtE)) \text{ AND } MemtoRegE$

$StallF = StallD = FlushE = lwstall$

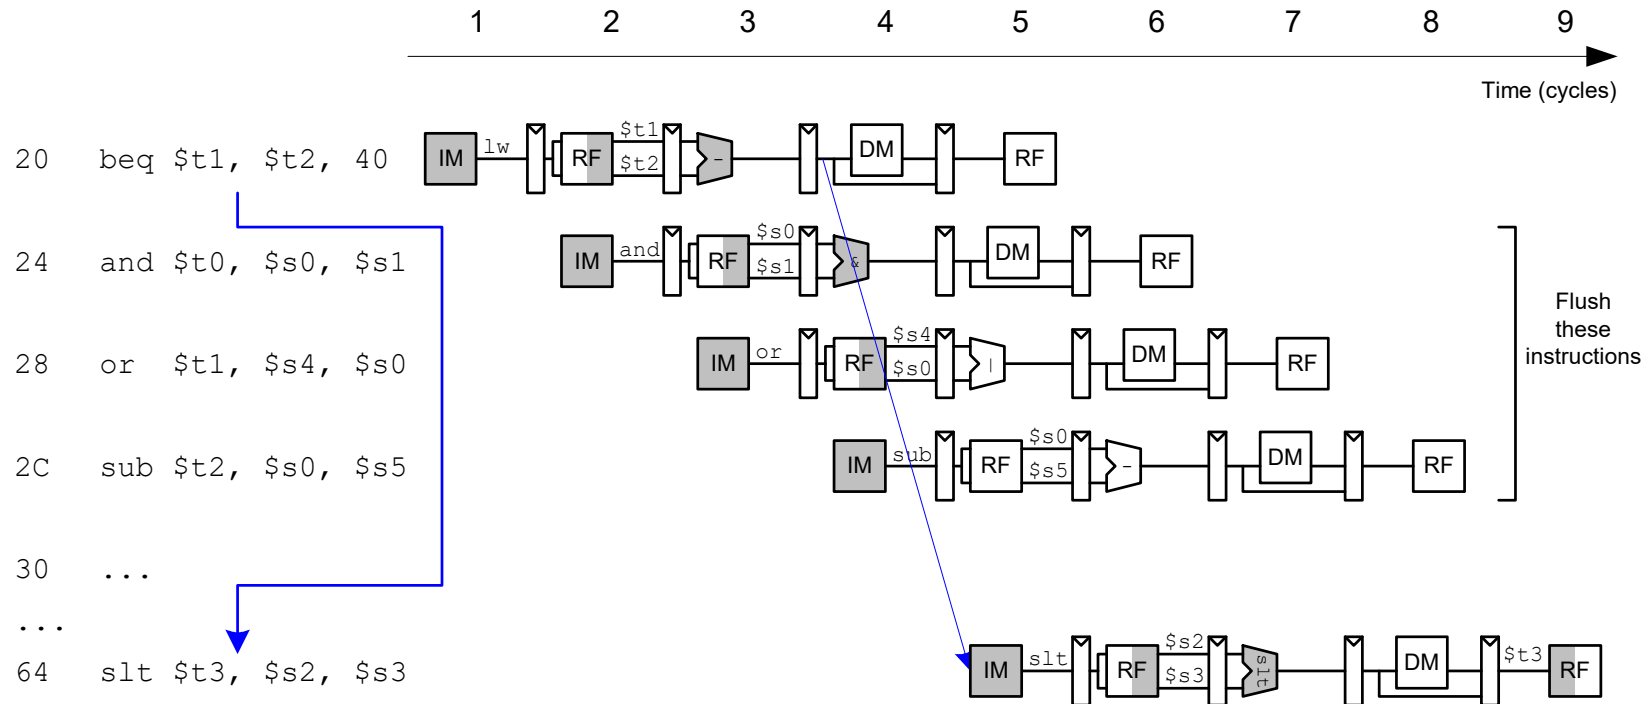
# Control Hazards

- **beq:**
  - branch not determined until 4<sup>th</sup> stage of pipeline
  - Instructions after branch fetched before branch occurs
  - These instructions must be flushed if branch happens
- **Branch misprediction penalty**
  - number of instruction flushed when branch is taken
  - May be reduced by determining branch earlier

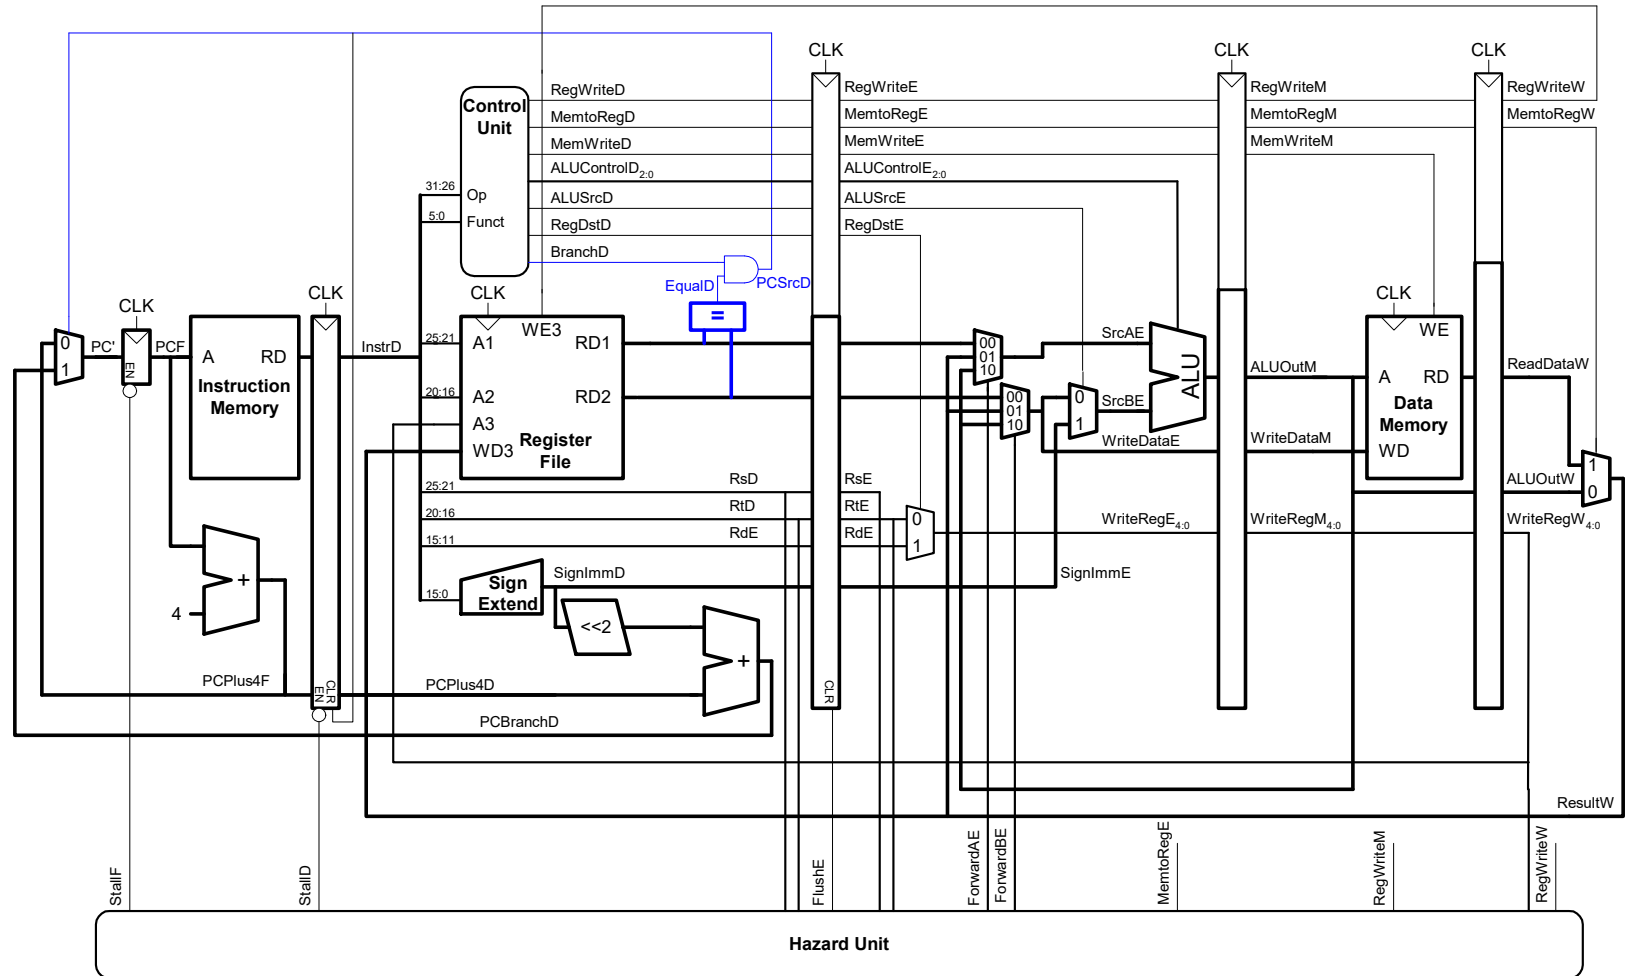
# Control Hazards: Original Pipeline



# Control Hazards

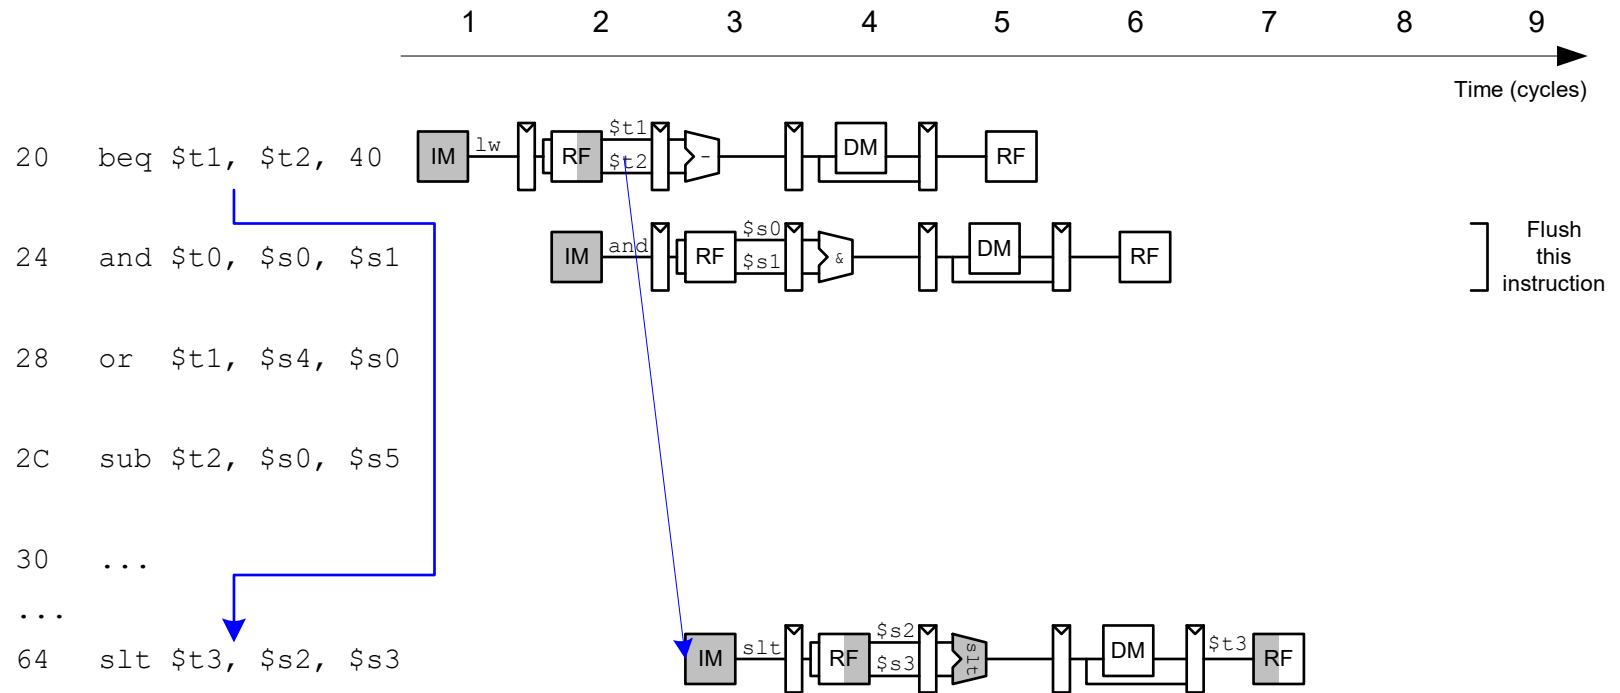


# Early Branch Resolution

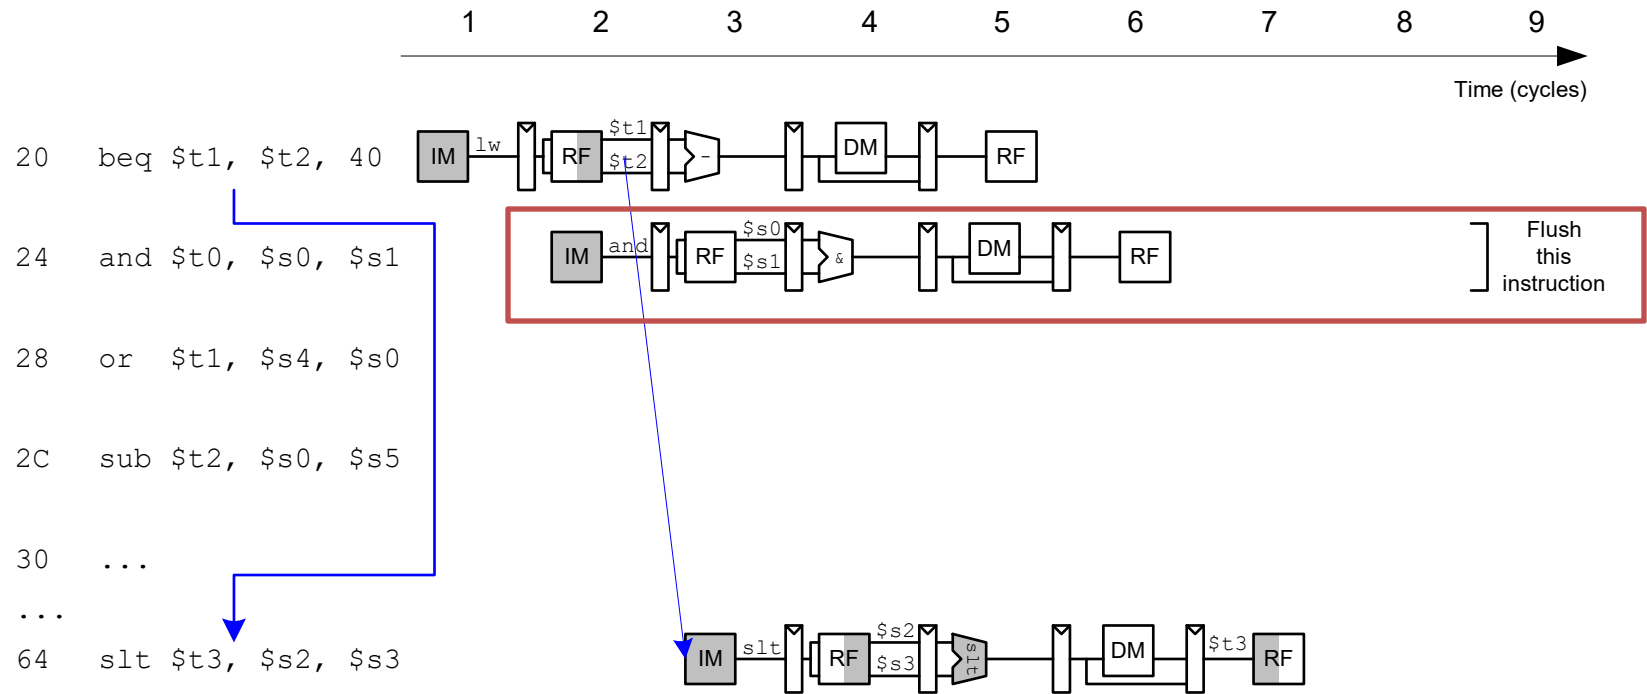


Introduced another data hazard in Decode stage

# Early Branch Resolution



# Branch delay slot

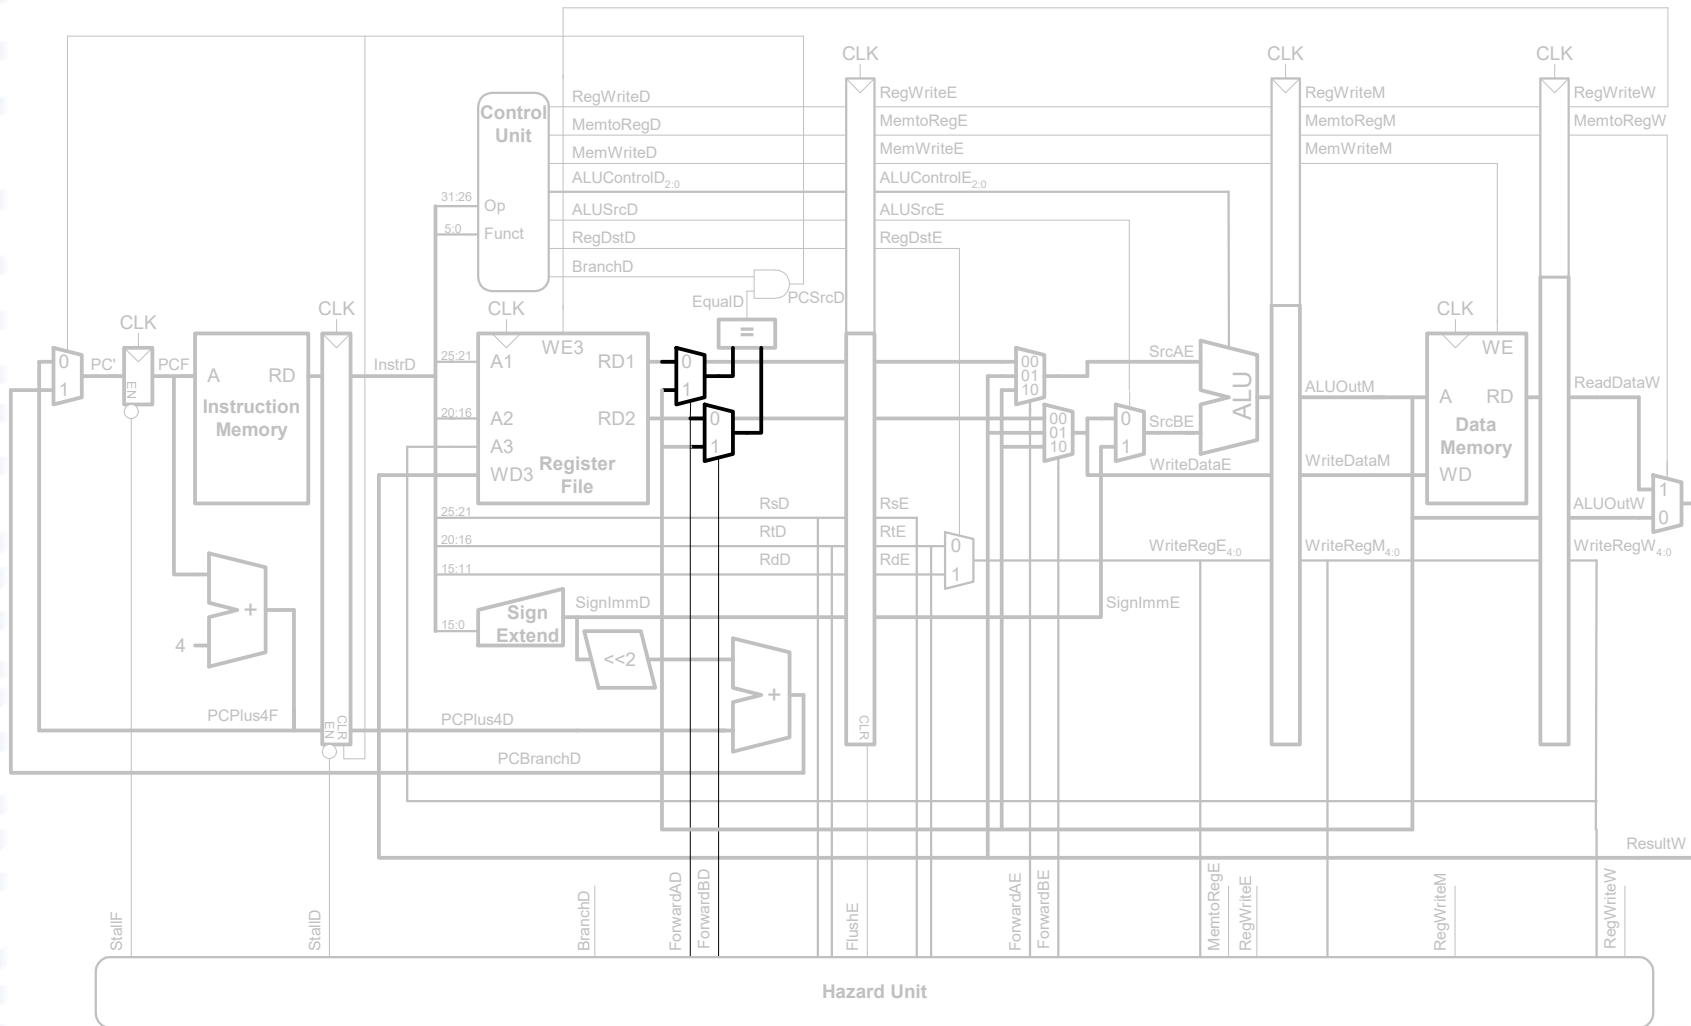


*Branch delay slot: The instruction after branch always execute.*

*In this way, the only left instruction could use effectively, instead of wasting this stage*



# Handling Data & Control Hazards



# Control Forwarding & Stalling Logic

- **Forwarding logic:**

**ForwardAD** = (*rsD* !=0) AND (*rsD* == WriteRegM) AND RegWriteM

**ForwardBD** = (*rtD* !=0) AND (*rtD* == WriteRegM) AND RegWriteM

- **Stalling logic:**

**branchstall** = BranchD AND RegWriteE AND  
(WriteRegE == *rsD* OR WriteRegE == *rtD*)

OR

BranchD AND MemtoRegM AND  
(WriteRegM == *rsD* OR WriteRegM == *rtD*)

**StallF** = **StallD** = **FlushE** = **lwstall** OR **branchstall**

# Branch Prediction

- Guess whether branch will be taken
  - Backward branches are usually taken (loops)
  - Consider history to improve guess
- Good prediction reduces fraction of branches requiring a flush

# Pipelined Performance Example

- SPECINT2000 benchmark:
  - 25% loads
  - 10% stores
  - 11% branches
  - 2% jumps
  - 52% R-type
- Suppose:
  - 40% of loads used by next instruction
  - 25% of branches mispredicted
  - All jumps flush next instruction
- **What is the average CPI?**

# Pipelined Performance Example

- SPECINT2000 benchmark:
  - 25% loads
  - 10% stores
  - 11% branches
  - 2% jumps
  - 52% R-type
- Suppose:
  - 40% of loads used by next instruction
  - 25% of branches mispredicted
  - All jumps flush next instruction
- **What is the average CPI?**
  - Load/Branch CPI = 1 when no stalling, 2 when stalling
  - $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$
  - $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$

$$\begin{aligned}\text{Average CPI} &= (0.25)(1.4) + (0.1)(1) + (0.11)(1.25) + (0.02)(2) + (0.52)(1) \\ &= \mathbf{1.15}\end{aligned}$$



# Pipelined Performance

- Pipelined processor critical path:

$$T_c = \max \{$$
$$t_{pcq} + t_{mem} + t_{setup}$$
$$2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup})$$
$$t_{pcq} + t_{mux} + t_{mux} + t_{ALU} + t_{setup}$$
$$t_{pcq} + t_{memwrite} + t_{setup}$$
$$2(t_{pcq} + t_{mux} + t_{RFwrite}) \}$$

# Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	30
Register setup	$t_{setup}$	20
Multiplexer	$t_{mux}$	25
ALU	$t_{ALU}$	200
Memory read	$t_{mem}$	250
Register file read	$t_{RFread}$	150
Register file setup	$t_{RFsetup}$	20
Equality comparator	$t_{eq}$	40
AND gate	$t_{AND}$	15
Memory write	$T_{memwrite}$	220
Register file write	$t_{RFwrite}$	100 ps

$$\begin{aligned} T_c &= 2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup}) \\ &= 2[150 + 25 + 40 + 15 + 25 + 20] \text{ ps} = \mathbf{550 \text{ ps}} \end{aligned}$$

# Pipelined Performance Example

Program with 100 billion instructions

$$\text{Execution Time} = (\# \text{ instructions}) \times \text{CPI} \times T_c$$

$$= (100 \times 10^9)(1.15)(550 \times 10^{-12})$$

$$= \mathbf{63 \text{ seconds}}$$



# Processor Performance Comparison

Processor	Execution Time (seconds)	Speedup (single-cycle as baseline)
<b>Single-cycle</b>	92.5	1
<b>Multicycle</b>	133	0.70
<b>Pipelined</b>	63	1.47