



Computer Architecture (Fall 2022)

Instruction-Level Parallelism

Dr. Duo Liu (刘铎)

Office: Main Building 0626

Email: liuduo@cqu.edu.cn

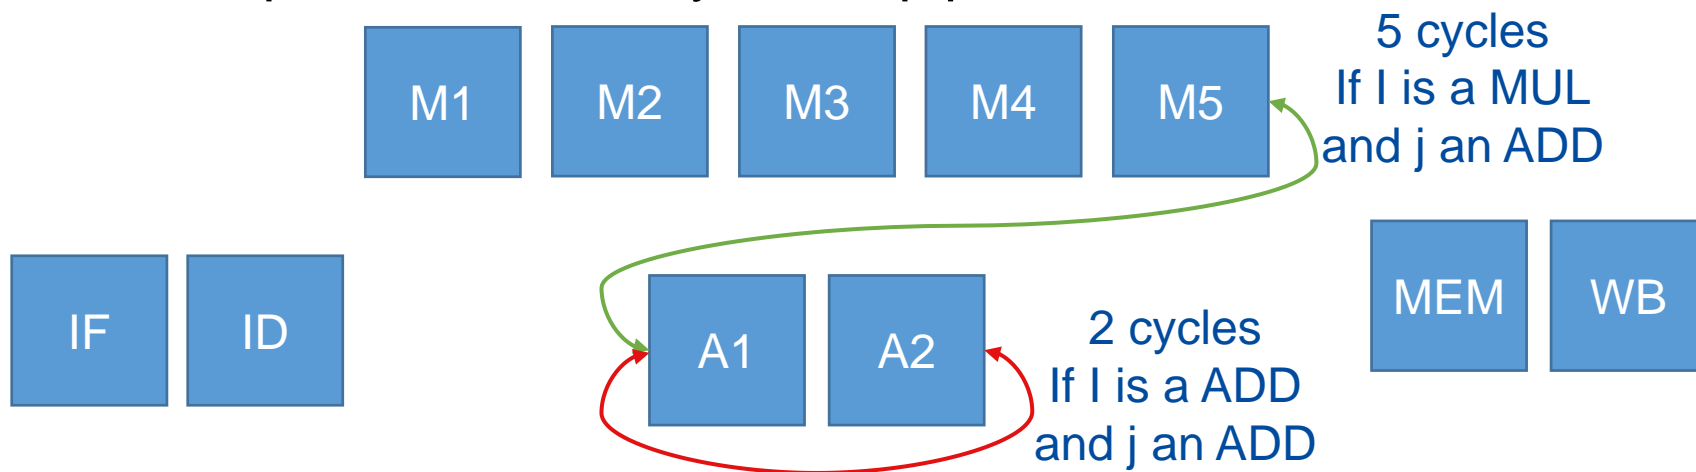
- Loop Unrolling & Pipeline Static Scheduling
- Software Pipelining
- Compiler-Based Branch Prediction

- Applications
 - processors that use static issue
 - VLIW and VLIW-like
 - NXP Semiconductors (founded by Philips) Trimedia
 - Transmeta Crusoe
 - processors that combine dynamic scheduling and static scheduling
 - Itanium Processor Family (based on IA-64 ISA)
 - Explicit Parallel Instruction Computing (EPIC)

How to Keep a Pipeline Full?

3/18

- Exploit Parallelism “in Time”
 - find sequences of independent instructions that can be overlapped in the pipeline
- Scheduling problem for a compiler
 - given “source instruction” *i* and “dependent instruction” *j*
 - avoid pipeline stalls by...
 - ...separating *j* from *i* by a distance in clock cycles that is equal to the latency of the pipeline execution of *i*



Example: a Simple Parallel Loop

4/18

```
for (i=1000; i>0; i=i-1) {  
    x[i] = x[i] + s;  
}
```



compiled into MIPS
assembly language

```
loop:  L.D      F0, 0(R1)      ; R1 = highest address of array element  
       ADD.D   F4, F0, F2     ; add scalar s (stored in F2)  
       S.D     F4, 0(R1)     ; store result  
       DADDUI  R1,R1,#-8      ; decrement pointer by 8 bytes (DW)  
       BNE     R1,R2, loop    ; 8(R2) points to x[1], last to process
```

Assumptions for Following Examples: Latencies and Resources of MIPS Pipeline

5/18

instruction Producing result	instruction Consuming result	Latency in Clock cycles
FP ALU op	FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0
Integer op	Integer op	1

- Additional assumption: no structural hazards
 - functional units are fully pipelined (or replicated as many times as the pipeline depth) so that an operation of any type can be issued on every clock cycle

Example: Scheduling the Simple Parallel Loop on the MIPS Pipeline

6/18

```
loop:  L.D      F0, 0(R1)    ; R1 = highest address of array element
      ADD.D    F4, F0, F2    ; add scalar s (which is stored in F2)
      S.D      F4, 0(R1)    ; store result
      DADDUI   R1,R1,#-8     ; decrement pointer by 8 bytes (DW)
      BNE     R1,R2, loop; 8(R2) points to x[1], last to process
```

• unscheduled execution

```
loop:  L.D      F0, 0(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      stall
      S.D      F4, 0(R1)
      DADDUI   R1, R1, #-8
      stall
      BNE     R1, R2, loop
```

• executionTime = 9 cycles

• scheduled execution

```
loop:  L.D      F0, 0(R1)
      DADDUI   R1, R1, #-8
      ADD.D    F4, F0, F2
      stall
      stall
      S.D      F4, 8(R1)
      BNE     R1, R2, loop
```

• executionTime = 7 cycles

• NOTE: execution time is expressed in 'clock cycles per loop iteration '

Scheduled Execution and Critical Path

7/18

- The clock cycle count for the loop is determined by the **critical path**, longest chain of dependent instructions

- in the example, 6 clock cycles to execute sequentially the chain

L.D F0, 0(R1)

stall

ADD.D F4, F0, F2

stall

stall

S.D F4, 8(R1)

- A good compiler (capable of some symbolic optimization) optimizes the “filling” of those stalls, for instance by altering and interchanging S.D with DADDUI

- **scheduled execution**

```
loop:  L.D      F0, 0(R1)
       DADDUI  R1, R1, #-8
       ADD.D   F4, F0, F2
       stall
       stall
       S.D     F4, 8(R1)
       BNE     R1, R2, loop
```

- executionTime = 7 cycles

Loop Overhead and Loop Unrolling

8/18

- At each iteration of the loop, the only actual work is done on the array

```
L.D      F0, 0(R1)
ADD.D    F4, F0, F2
S.D      F4, 8(R1)
```

- Beside the stall cycle, **DADDUI** and **BNE** represent **loop overhead** cycles
- To reduce loop overhead, before scheduling the loop execution apply **loop unrolling**
 - replicate the loop body multiple times
 - merge unnecessary instructions
 - eliminate the name dependencies
 - adjust the control code

- scheduled execution**

```
loop:  L.D      F0, 0(R1)
       DADDUI   R1, R1, #-8
       ADD.D    F4, F0, F2
       stall
       stall
       S.D      F4, 8(R1)
       BNE      R1, R2, loop
```

- executionTime = 7 cycles**

Reducing Loop Overhead by Loop Unrolling before Scheduling – Step 1 (replication)

9/18

- original loop

```
loop:  L.D    F0, 0(R1)
       ADD.D  F4, F0, F2
       S.D    F4, 0(R1)
       DADDUI R1, R1, #-8
       BNE    R1, R2, loop
```



- unrolling the original loop 4 times (dropping 3 **BNE**)

```
loop:  L.D    F0, 0(R1)
       ADD.D  F4, F0, F2
       S.D    F4, 0(R1)
       DADDUI R1, R1, #-8
       L.D    F0, 0(R1)
       ADD.D  F4, F0, F2
       S.D    F4, 0(R1)
       DADDUI R1, R1, #-8
       L.D    F0, 0(R1)
       ADD.D  F4, F0, F2
       S.D    F4, 0(R1)
       DADDUI R1, R1, #-8
       L.D    F0, 0(R1)
       ADD.D  F4, F0, F2
       S.D    F4, 0(R1)
       DADDUI R1, R1, #-8
       BNE    R1, R2, loop
```

Execution Time after Step 1

10/18

- unrolling the original loop 4 times (dropping 3 **BNE**)

```

loop:  L.D      F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        L.D     F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        L.D     F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        L.D     F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        BNE     R1, R2, loop
    
```

```

loop:  L.D      F0, 0(R1)
        stall
        ADD.D   F4, F0, F2
        stall
        stall
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        stall
        L.D     F0, 0(R1)
        stall
        ADD.D   F4, F0, F2
        stall
        stall
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        stall
        L.D     F0, 0(R1)
        stall
        ADD.D   F4, F0, F2
        stall
        stall
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        stall
        BNE     R1, R2, loop
    
```

- executionTime = 33 cycles

Reducing Loop Overhead by Loop Unrolling before Scheduling – Step 2 (merging)

11/18

- unrolling the original loop 4 times (dropping 3 **BNE**)

```
loop:  L.D      F0, 0(R1)
      ADD.D   F4, F0, F2
      S.D     F4, 0(R1)
      DADDUI  R1, R1, #-8
      L.D     F0, 0(R1)
      ADD.D   F4, F0, F2
      S.D     F4, 0(R1)
      DADDUI  R1, R1, #-8
      L.D     F0, 0(R1)
      ADD.D   F4, F0, F2
      S.D     F4, 0(R1)
      DADDUI  R1, R1, #-8
      L.D     F0, 0(R1)
      ADD.D   F4, F0, F2
      S.D     F4, 0(R1)
      DADDUI  R1, R1, #-8
      BNE     R1, R2, loop
```



- merging **DADDUI** and adjusting the offsets

```
loop:  L.D      F0, 0(R1)
      ADD.D   F4, F0, F2
      S.D     F4, 0(R1)
      L.D     F0, -8(R1)
      ADD.D   F4, F0, F2
      S.D     F4, -8(R1)
      L.D     F0, -16(R1)
      ADD.D   F4, F0, F2
      S.D     F4, -16(R1)
      L.D     F0, -24(R1)
      ADD.D   F4, F0, F2
      S.D     F4, -24(R1)
      DADDUI  R1, R1, #-32
      BNE     R1, R2, loop
```

Execution Time after Step 2

12/18

- merging **DADDUI** and adjusting the offsets

```
loop:  L.D      F0, 0(R1)
      ADD.D    F4, F0, F2
      S.D      F4, 0(R1)
      L.D      F0, -8(R1)
      ADD.D    F4, F0, F2
      S.D      F4, -8(R1)
      L.D      F0, -16(R1)
      ADD.D    F4, F0, F2
      S.D      F4, -16(R1)
      L.D      F0, -24(R1)
      ADD.D    F4, F0, F2
      S.D      F4, -24(R1)
      DADDUI   R1, R1, #-32
      BNE     R1, R2, loop
```

```
loop:  L.D      F0, 0(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      S.D      F4, 0(R1)
      L.D      F0, -8(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      S.D      F4, -8(R1)
      L.D      F0, -16(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      S.D      F4, -16(R1)
      L.D      F0, -24(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      S.D      F4, -24(R1)
      DADDUI   R1, R1, #-32
      stall
      BNE     R1, R2, loop
```

- executionTime = 27 cycles

Reducing Loop Overhead by Loop Unrolling before Scheduling – Step 3 (renaming)

13/18

- merging **DADDUI** and adjusting the offsets

```
loop:  L.D      F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        L.D     F0, -8(R1)
        ADD.D   F4, F0, F2
        S.D     F4, -8(R1)
        L.D     F0, -16(R1)
        ADD.D   F4, F0, F2
        S.D     F4, -16(R1)
        L.D     F0, -24(R1)
        ADD.D   F4, F0, F2
        S.D     F4, -24(R1)
        DADDUI  R1, R1, #-32
        BNE     R1, R2, loop
```



- eliminating name dependences via register renaming

```
loop:  L.D      F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        L.D     F6, -8(R1)
        ADD.D   F8, F6, F2
        S.D     F8, -8(R1)
        L.D     F10, -16(R1)
        ADD.D   F12, F10, F2
        S.D     F12, -16(R1)
        L.D     F14, -24(R1)
        ADD.D   F16, F14, F2
        S.D     F16, -24(R1)
        DADDUI  R1, R1, #-32
        BNE     R1, R2, loop
```

Execution Time after Step 3

14/18

- eliminating name dependences via register renaming

```
loop:  L.D      F0, 0(R1)
      ADD.D    F4, F0, F2
      S.D      F4, 0(R1)
      L.D      F6, -8(R1)
      ADD.D    F8, F6, F2
      S.D      F8, -8(R1)
      L.D      F10, -16(R1)
      ADD.D    F12, F10, F2
      S.D      F12, -16(R1)
      L.D      F14, -24(R1)
      ADD.D    F16, F14, F2
      S.D      F16, -24(R1)
      DADDUI   R1, R1, #-32
      BNE     R1, R2, loop
```

```
loop:  L.D      F0, 0(R1)
      stall
      ADD.D    F4, F0, F2
      stall
      stall
      S.D      F4, 0(R1)
      L.D      F6, -8(R1)
      stall
      ADD.D    F8, F6, F2
      stall
      stall
      S.D      F8, -8(R1)
      L.D      F10, -16(R1)
      stall
      ADD.D    F12, F10, F2
      stall
      stall
      S.D      F12, -16(R1)
      L.D      F14, -24(R1)
      stall
      ADD.D    F16, F14, F2
      stall
      stall
      S.D      F16, -24(R1)
      DADDUI   R1, R1, #-32
      stall
      BNE     R1, R2, loop
```

- executionTime = 27 cycles

Reducing Loop Overhead by Loop Unrolling before Scheduling – Step 4 (scheduling)

15/18

- eliminating name dependences via register renaming

```
loop:  L.D      F0, 0(R1)
       ADD.D    F4, F0, F2
       S.D      F4, 0(R1)
       L.D      F6, -8(R1)
       ADD.D    F8, F6, F2
       S.D      F8, -8(R1)
       L.D      F10, -16(R1)
       ADD.D    F12, F10, F2
       S.D      F12, -16(R1)
       L.D      F14, -24(R1)
       ADD.D    F16, F14, F2
       S.D      F16, -24(R1)
       DADDUI   R1, R1, #-32
       BNE     R1, R2, loop
```



- scheduling

```
loop:  L.D      F0, 0(R1)
       L.D      F6, -8(R1)
       L.D      F10, -16(R1)
       L.D      F14, -24(R1)
       ADD.D    F4, F0, F2
       ADD.D    F8, F6, F2
       ADD.D    F12, F10, F2
       ADD.D    F16, F14, F2
       S.D      F4, 0(R1)
       S.D      F8, -8(R1)
       DADDUI   R1, R1, #-32
       S.D      F12, 16(R1)
       S.D      F16, 8(R1)
       BNE     R1, R2, loop
```

Scheduling after Step 4

16/18

- scheduling

```
loop:  L.D      F0, 0(R1)
       L.D      F6, -8(R1)
       L.D      F10, -16(R1)
       L.D      F14, -24(R1)
       ADD.D    F4, F0, F2
       ADD.D    F8, F6, F2
       ADD.D    F12, F10, F2
       ADD.D    F16, F14, F2
       S.D      F4, 0(R1)
       S.D      F8, -8(R1)
       DADDUI   R1, R1, #-32
       S.D      F12, 16(R1)
       BNE     R1, R2, loop
       S.D      F16, 8(R1)
```

```
loop:  L.D      F0, 0(R1)
       L.D      F6, -8(R1)
       L.D      F10, -16(R1)
       L.D      F14, -24(R1)
       ADD.D    F4, F0, F2
       ADD.D    F8, F6, F2
       ADD.D    F12, F10, F2
       ADD.D    F16, F14, F2
       S.D      F4, 0(R1)
       S.D      F8, -8(R1)
       DADDUI   R1, R1, #-32
       S.D      F12, 16(R1)
       S.D      F16, 8(R1)
       BNE     R1, R2, loop
```

- executionTime = 14 cycles

Final Performance Comparison

17/18

- original scheduled execution

```
Loop:  L.D      F0, 0(R1)
        DADDUI  R1, R1, #-8
        ADD.D   F4, F0, F2
        stall
        stall
        S.D     F4, 8(R1)
        BNE     R1, R2, loop
```

- executionTime = 7cycles
- execution Time Per Element = 7 cycles
- e.g., an array of 1000 elements is processed in 7,000 cycles

- scheduled execution after loop unrolling

```
loop:   L.D      F0, 0(R1)
        L.D      F6, -8(R1)
        L.D      F10, -16(R1)
        L.D      F14, -24(R1)
        ADD.D    F4, F0, F2
        ADD.D    F8, F6, F2
        ADD.D    F12, F10, F2
        ADD.D    F16, F14, F2
        S.D      F4, 0(R1)
        S.D      F8, -8(R1)
        DADDUI   R1, R1, #-32
        S.D      F12, 16(R1)
        S.D      F16, 8(R1)
        BNE     R1, R2, loop
```

- executionTime = 14 cycles
- execution Time Per Element = 3.5 cycles
- e.g., an array of 1000 elements is processed in 3,500 cycles

Loop-unrolling speedup = 2

- How is the final instruction sequence obtained?
- The compiler needs to determine that it is both possible and convenient to...
 - **unroll the loop** because the iterations are independent (except for the loop maintenance code)
 - **eliminate extra loop maintenance code** as long as the offsets are properly adjusted
 - **use different registers** to increase parallelism by removing name dependencies
 - **remove stalling cycles** by interleaving the instructions (and adjusting the offset or index increment consequently)
 - e.g. the compiler analyzes the memory addresses to determine that loads and stores from different iterations are independent