

Computer Architecture (Spring 2020)

Instruction Set Principles

Dr. Duo Liu (刘铎)
Office: Main Building 0626
Email: liuduo@cqu.edu.cn

MIPS


- **1981:** Stanford MIPS Computer, a “**M**icroprocessor without **I**nterlocked **P**ipeline **S**tages” [Hennessy81]
 - no HW to stall the pipeline (handling dependencies is compiler’s job)
- **1984:** Hennessy founds MIPS Computer System
 - R2000 & R3000 first products (with interlock in HW!)
- **1991:** MIPS releases the 64-bit R4000
 - SGI buys MIPS which becomes the division MIPS Technologies
- **1999:** MIPS-Technologies so successful that SGI spins off it
 - two products MIPS32, MIPS64
 - major revenues portion from licensing the design
 - estimated 1/3 of the produced RISCs is a MIPS-based design
 - almost 100 million of MIPS manufactured in 2002
 - Used in products from ATI Technologies, Broadcom, Cisco, NEC, Nintendo
SGI, Sony, Texas Instruments, Toshiba...
- **MIPS64** is used throughout the textbook to illustrate the ILP techniques

Before Introducing MIPS64: What We Expect?

- Architecture?
 - register-to-register (load-store)
- Registers?
 - “many” general purpose registers and dedicated FP registers
- Addressing modes support?
 - displacement (address-offset of 12 to 16 bits)
 - using PC-relative addressing, branch +/- 2^{15} words from PC
 - immediate (size 8-16 bits)
 - register indirect
- Data sizes & types?
 - Integers (8,16,32,64 size) and FP (64)
- Instructions types?
 - strong support for simple pervasive instructions
 - load, store, add, sub, move register-register, and shift
 - compare equal/not equal/less, branch, jump, call, return
- Encoding?
 - fixed for performance rather than variable for code size

MIPS - Registers

- 32 64-bit GPRs
 - R0 is hardwired to zero
 - AT, GP, SP, FP, RA are conventionally mapped to R1, R28, R29, R30, R31
- 32 Floating Point Registers holding either
 - 32 single-precision values (32 bits)
 - half of the FPR is not used
 - 32 double-precision values (64 bits)
- Special registers
 - PC – program counter
 - HI, LO
 - integer multiply result (higher and lower word)
 - integer divide result (remainder and quotient)



MIPS – Data Types

- Integer data
 - 8-bit **bytes**
 - 16-bit **half words**
 - present in C and popular in programs like operating systems (concerned about size of data structures)
 - also even more popular if Unicode becomes more pervasive
 - 32-bit **words**
 - 64-bit **double words**
- Floating Point
 - 32-bit **single precision**
 - Same motivations as for the 16-bit half words
 - 64-bit **double precision**

These are loaded onto the GPRs with either additional zeros or the sign bit and operated upon with 64-bit integer operations

MIPS64 operations work directly on these

MIPS – Addressing Modes

- The hardware only support two addressing modes

- **immediate** (16-bit)

DADDI R3, R2, #7 ; $\text{Regs}[\text{R3}] \leftarrow \text{Regs}[\text{R2}] + 7$

- **displacement** (16-bit)

LD R3, 100(R1) ; $\text{Regs}[\text{R3}] \leftarrow \text{Mem}[100 + \text{Regs}[\text{R1}]]$

- Two other modes supported “indirectly”

- **register indirect** (placing 0 in the 16-bit displacement field)

LD R3, 0(R1) ; $\text{Regs}[\text{R3}] \leftarrow \text{Mem}[\text{Regs}[\text{R1}]]$

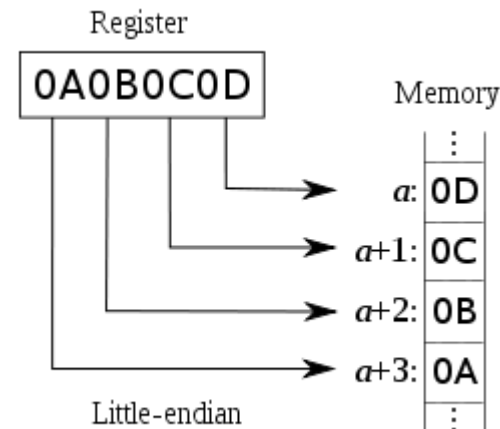

- **direct addressing** (using 0 as the base register)

LD R3, 1001(R0) ; $\text{Regs}[\text{R3}] \leftarrow \text{Mem}[1001]$

- Memory is byte-addressable with a 64-bit address

Interpreting Memory Addresses

- Memory word is addressed by the byte at the lowest address
 - **Big Endian**
 - Most significant byte is at the lowest address
 - “big end first”
 - Ex: Motorola 68000, SPARC, IBM 360
 - **Little Endian**
 - Least significant byte is at the lowest address
 - “little end first”
 - Ex: Intel x86, AMD64, VAX



Some architectures can be configured in both ways: ARM, MIPS, IA64, PowerPC

MIPS – Instruction Format and Layout

- MIPS instructions fall into 5 classes:
 - Arithmetic/logical/shift/comparison
 - Control instructions (branch and jump)
 - Load/store
 - Other (exception, register movement to/from GP registers, etc.)
- Three instruction encoding formats:
 - R-type (6-bit opcode, 5-bit rs, 5-bit rt, 5-bit rd, 5-bit shamt, 6-bit function code)

31-26	25-21	20-16	15-11	10-6	5-0
<i>opcode</i>	<i>rs</i>	<i>rt</i>	<i>rd</i>	<i>shamt</i>	<i>function</i>

- I-type (6-bit opcode, 5-bit rs, 5-bit rt, 16-bit immediate)

31-26	25-21	20-16	15-0
<i>opcode</i>	<i>rs</i>	<i>rt</i>	<i>imm</i>

- J-type (6-bit opcode, 26-bit pseudo-direct address)

31-26	25-0
<i>opcode</i>	<i>pseudodirect jump address</i>

Example Instructions

- **ADD \$2, \$3, \$4**
 - R-type A/L/S/C instruction
 - Opcode is 0's, rd=2, rs=3, rt=4, func=000010
 - 000000 00011 00100 00010 00000 000010
- **JALR \$3**
 - R-type jump instruction
 - Opcode is 0's, rs=3, rt=0, rd=31 (by default), func=001001
 - 000000 00011 00000 11111 00000 001001
- **ADDI \$2, \$3, 12**
 - I-type A/L/S/C instruction
 - Opcode is 001000, rs=3, rt=2, imm=12
 - 001000 00011 00010 00000000000001100

Example Instructions

- BEQ \$3, \$4, 4
 - I-type conditional branch instruction
 - Opcode is 000100, rs=00011, rt=00100, imm=4 (skips next 4 instructions)
 - 000100 00011 00100 0000000000000000100
- SW \$2, 128(\$3)
 - I-type memory address instruction
 - Opcode is 101011, rs=00011, rt=00010, imm=0000000010000000
 - 101011 00011 00010 0000000010000000
- J 128
 - J-type pseudodirect jump instruction
 - Opcode is 000010, 26-bit pseudodirect address is $128/4 = 32$
 - 000010 0000000000000000000000100000

- ```

;Mem[40 + Regs[R3] ←32 Regs[F0]]0..31

```

# MIPS – Some Arithmetic/Logic Instructions

---

- Double Word Add

**DADD R1, R2, R3** ;Regs[R1]  $\leftarrow$  Regs[R2] + Regs[R3]

- Double Word Add Immediate Unsigned

**DADDIU R1, R2, #3** ;Regs[R1]  $\leftarrow$  Regs[R2] + 3

- Double Word Shift Left Logical

**DSLL R1, R2, #5** ;Regs[R1]  $\leftarrow$  Regs[R2]  $\ll$  5

- Set on Less Than

**SLT R1, R2, R3** ; if (Regs[R2] < Regs[R3])

Regs[R1]  $\leftarrow$  1 else Regs[R1]  $\leftarrow$  0

- Multiply and Add Word to HI,LO registers

**MADD R1, R2** ; (LO,HI)  $\leftarrow$  Regs[R1] x Regs[R2] + (LO,HI)

- Loading a Constant (mnemonic: **LI**)

**LI R1, #3** ;Regs[R1]  $\leftarrow$  3

- Register-Register Move (mnemonic: **MOV**)

**MOV R1, R2** ;Regs[R1]  $\leftarrow$  Regs[R2]

## EXAMPLE

---

```
Loop: add $t1, $s3, $s3 # starts from 80000
 add $t1, $t1, $t1
 add $t1, $t1, $s6
 lw $t0, 0($t1)
 bne $t0, $s5, Exit
 add $s3, $s3, $s4
 j Loop
```

Exit:

# EXAMPLE

---

|       | 6  | 5     | 5  | 5  | 5 | 6  |        |
|-------|----|-------|----|----|---|----|--------|
| 80000 | 0  | 19    | 19 | 9  | 0 | 32 | R-type |
| 80004 | 0  | 9     | 9  | 9  | 0 | 32 | R-type |
| 80008 | 0  | 9     | 22 | 9  | 0 | 32 | R-type |
| 80012 | 35 | 9     | 8  | 0  |   |    | I-type |
| 80016 | 5  | 8     | 21 | 2  |   |    | I-type |
| 80020 | 0  | 19    | 20 | 19 | 0 | 32 | R-type |
| 80024 | 2  | 20000 |    |    |   |    | J-type |
| 80028 |    |       |    |    |   |    |        |

# MIPS – Some Arithmetic/Logic Instructions

---

- Suppose register R3 has the binary number

1111 1111 1111 1111 1111 1111 1111 1111

- and register R4 has the binary number

0000 0000 0000 0000 0000 0000 0000 0001

- What are the values of R1 and R2 after this?

**SLT R1 R3 R4**

**SLTU R2 R3 R4**

- Answer:

**R1 = 1, R2 = 0**

- because

R3 = -1 (as an integer) and =4,294,967,295 (as unsigned)

R4 = 1 in both cases

# MIPS – Some Control-Flow Instructions

---

- **Jump** (28-bit target address is 26-bit name shifted left twice)

**J name** ;  $PC \leftarrow PC_{64..28} \text{ ## name ## } 0^2$

- **Jump and link** (using the return address register  $RA \equiv R31$ )

**JAL name** ;  $\text{Regs}[R31] \leftarrow PC+8$ ;  $PC \leftarrow PC_{64..28} \text{ ## name ## } 0^2$

- **Jump and link register**

**JALR R2** ;  $\text{Regs}[R31] \leftarrow PC+8$ ;  $PC \leftarrow \text{Regs}[R2]$

- **Jump register**

**JR R3** ;  $PC \leftarrow \text{Regs}[R3]$

- **Branch equal (zero)** (compilers combined this with SLT, SLTI in order to build branch-on-less conditions)

**BEQ R4, R0, name** ; if ( $\text{Regs}[R4] == 0$ )  $PC \leftarrow PC + \text{signExt}\{\text{name ## } 0^2\}$

- **Branch not equal**

**BNE R3, R4, name** ; if ( $\text{Regs}[R3] \neq \text{Regs}[R4]$ )  $PC \leftarrow PC + \text{signExt}\{\text{name ## } 0^2\}$

- **Conditional move if zero** (support conversion of a simple branch into a conditional arithmetic instruction)

**MOVZ R1, R2, R3** ; if ( $\text{Regs}[R3] == 0$ )  $\text{Regs}[R1] \leftarrow \text{Regs}[R2]$



# Review: MIPS-Jump Instruction Example

```
/* Assume loop stored at
 0000 A000 6BC4 0308 */
Loop: SLL R9,R19,2
 ADD R9,R9,R22
 LW R8, 0(R9)
 BNE R8, R21,exit
 ADDI R19,R19,1
 J Loop
Exit: LW R11, 0(R18)
```

- j loop : [PC-region addressing] the target address effectively encoded is 2F9 00C2 that shifted-left by two and composed with the upper 36 bits of the PC of the following instruction gives 0000 a000 6bc4 0308
- BNE R8, R21, exit : [PC-relative addressing] the offset is 2 (it gets summed to the PC of the following instruction to point to Exit)

|                       |      |                           |     |    |        |     |
|-----------------------|------|---------------------------|-----|----|--------|-----|
| 0000 a000 6bc4 0308 : | SLL  | 0                         | R19 | R9 | 2      | sll |
| 0000 a000 6bc4 030C : | ADD  | R9                        | R22 | R9 | 000000 | add |
| 0000 a000 6bc4 0310 : | LW   | R9                        | R8  | 0  |        |     |
| 0000 a000 6bc4 0314 : | BNE  | R8                        | R21 | 2  |        |     |
| 0000 a000 6bc4 0318 : | ADDI | R19                       | R19 | 1  |        |     |
| 0000 a000 6bc4 031C : | JUMP | 2F9 00C2 = (bc4 0308 / 4) |     |    |        |     |
| 0000 a000 6bc4 0320 : | LW   | R18                       | R11 | 0  |        |     |

# Review: MIPS-Jump Instruction Example

```
/* Assume loop stored at
0000 A000 6BC4 0308 */
Loop: SLL R9,R19,2
 ADD R9,R9,R22
 LW R8, 0(R9)
 BNE R8, R21,exit
 ADDI R18,R18,1
 J Loop
Exit: LW R11, 0(R18)
```

- `j loop` : [PC-region addressing] the target address effectively encoded is `2F9 00C2` that shifted-left by two and composed with the upper 36 bits of the PC of the following instruction gives `0000 a000 6bc4 0308`
- `BNE R8, R21, exit` : [PC-relative addressing] the offset is 2 (it gets summed to the PC of the following instruction to point to Exit)

|                       |        |                                  |        |        |        |        |
|-----------------------|--------|----------------------------------|--------|--------|--------|--------|
| 0000 a000 6bc4 0308 : | 000000 | 000000                           | 010011 | 001001 | 000010 | 000000 |
| 0000 a000 6bc4 030C : | 000000 | 001001                           | 010110 | 001001 | 000000 | 100000 |
| 0000 a000 6bc4 0310 : | 100011 | 001001                           | 001000 | 000000 |        |        |
| 0000 a000 6bc4 0314 : | 000101 | 001000                           | 010101 | 000010 |        |        |
| 0000 a000 6bc4 0318 : | 001000 | 010011                           | 010011 | 000001 |        |        |
| 0000 a000 6bc4 031C : | 000010 | 1011 1110 0100 0000 0011 0000 10 |        |        |        |        |
| 0000 a000 6bc4 0320 : | 100011 | 010010                           | 001011 | 000000 |        |        |

# MIPS – Floating-Point Instructions

---

- Floating Point Absolute Value (Double Precision)

**ABS.D F1, F2** ;F1  $\leftarrow$  abs(F2)

- Floating Point Add (Single Precision)

**ADD.S F3, F1, F2** ;Regs[F3]  $\leftarrow$  Regs[F1] + Regs[F2]

- Floating Point Add (Double Precision)

**ADD.D F3, F1, F2** ;Regs[F3]  $\leftarrow$  Regs[F1] + Regs[F2]

- Floating Point Divide (Double Precision)

**DIV.D F3, F1, F2** ;Regs[F3]  $\leftarrow$  Regs[F1] / Regs[F2]

- Floating Point Square Root (Double Precision)

**SQRT.D F3, F1** ;Regs[F3]  $\leftarrow$  sqrt(Regs[F1])

- Floating Point Multiply Add (Single Precision)

**MADD.S F3, F1, F2, F4** ;Regs[F3]  $\leftarrow$  (Regs[F2] x Regs[F4]) +  
+ Regs[F1]

# Real World Instruction Sets

| Arch      | Type    | # Oper | # Mem | Data Size      | # Regs | Addr Size | Use                   |
|-----------|---------|--------|-------|----------------|--------|-----------|-----------------------|
| Alpha     | Reg-Reg | 3      | 0     | 64-bit         | 32     | 64-bit    | Workstation           |
| ARM       | Reg-Reg | 3      | 0     | 32/64-bit      | 16     | 32/64-bit | Cell Phones, Embedded |
| MIPS      | Reg-Reg | 3      | 0     | 32/64-bit      | 32     | 32/64-bit | Workstation, Embedded |
| SPARC     | Reg-Reg | 3      | 0     | 32/64-bit      | 24-32  | 32/64-bit | Workstation           |
| TI C6000  | Reg-Reg | 3      | 0     | 32-bit         | 32     | 32-bit    | DSP                   |
| IBM 360   | Reg-Mem | 2      | 1     | 32-bit         | 16     | 24/31/64  | Mainframe             |
| x86       | Reg-Mem | 2      | 1     | 8/16/32/64-bit | 4/8/24 | 16/32/64  | Personal Computers    |
| VAX       | Mem-Mem | 3      | 3     | 32-bit         | 16     | 32-bit    | Minicomputer          |
| Mot. 6800 | Accum.  | 1      | 1/2   | 8-bit          | 0      | 16-bit    | Microcontroler        |

# Why the Diversity in ISAs?

---

## Technology Influenced ISA

- Storage is expensive, tight encoding important
- Reduced Instruction Set Computer
  - Remove instructions until whole computer fits on die
- Multicore/Manycore
  - Transistors not turning into sequential performance

## Application Influenced ISA

- Instructions for Applications
  - DSP instructions
- Compiler Technology has improved
  - SPARC Register Windows no longer needed
  - Compiler can register allocate effectively

# Quiz 2

1. H&P5th Page A-47 A.1

2.


[Self-Study]:

The value represented **0x(0123 4567 89AB CDEF)** is to be stored in an aligned **64bit** double word. Choose an address from **0x00002 to 0x00012** as the starting address, and store them using Little Endian byte order.

# Assigned Readings

---

- Computer Architecture – A Quantitative Approach by  
**John Hennessy**
  - Stanford University**Dave Patterson**
  - UC BerkeleyFifth Edition - 2012  
Morgan Kaufmann (Elsevier)
- Read Sections  
Appendix A.1-A.7, A.9-A.11



# Homework #2

---

H&P 5th

Page A-47: A.3, A.7, A.9

Due 2020/04/09